

Python for scientific computing

Python has extensive packages to help with data analysis:

- numpy: matrices, linear algebra, Fourier transform, pseudorandom number generators
- scipy: advanced linear algebra and maths, signal processing, statistics
- pandas: DataFrames, data wrangling and analysis
- matplotlib: visualizations such as line charts, histograms, scatter plots.

<IPython.core.display.HTML object>

NumPy

NumPy is the fundamental package required for high performance scientific computing in Python. It provides:

- ndarray: fast and space-efficient n-dimensional numeric array with vectorized arithmetic operations
- Functions for fast operations on arrays without having to write loops
- Linear algebra, random number generation, Fourier transform
- Integrating code written in C, C++, and Fortran (for faster operations)

pandas provides a richer, simpler interface to many operations. We'll focus on using ndarrays here because they are heavily used in scikit-learn.

ndarrays

There are several ways to create numpy arrays.

```
[2]: # Convert normal Python array to 1-dimensional numpy array
      np.array((1, 2, 53))
```

```
array([ 1,  2, 53])
```

```
[3]: # Convert sequences of sequences of sequences ... to n-dim array
      np.array([(1.5, 2, 3), (4, 5, 6)])
```

```
array([[ 1.5,  2. ,  3. ],
       [ 4. ,  5. ,  6. ]])
```

```
[4]: # Define element type at creation time
      np.array([[1, 2], [3, 4]], dtype=complex)
```

```
array([[ 1.+0.j,  2.+0.j],
       [ 3.+0.j,  4.+0.j]])
```

Useful properties of ndarrays:

```
[5]: my_array = np.array([[1, 0, 3], [0, 1, 2]])
      my_array.ndim      # number of dimensions (axes), also called the rank
      my_array.shape     # a matrix with n rows and m columns has shape (n,m)
      my_array.size      # the total number of elements of the array
      my_array.dtype     # type of the elements in the array
      my_array.itemsize  # the size in bytes of each element of the array
```

2

(2, 3)

6

dtype('int64')

8

Quick array creation.

It is cheaper to create an array with placeholders than extending it later.

```
[6]: np.ones(3) # Default type is float64
      np.zeros([2, 2])
      np.empty([2, 2]) # Fills the array with whatever sits in memory
      np.random.random((2,3))
      np.random.randint(5, size=(2, 4))
```

```
array([ 1.,  1.,  1.])
```

```
array([[ 0.,  0.],
       [ 0.,  0.]])
```

```
array([[ 0.,  0.],
       [ 0.,  0.]])
```

```
array([[ 0.404,  0.114,  0.042],
       [ 0.866,  0.599,  0.222]])
```

```
array([[4, 3, 2, 0],
       [2, 0, 2, 3]])
```

Create sequences of numbers

```
[7]: np.linspace(0, 1, num=4) # Linearly distributed numbers between 0 and 1
      np.arange(0, 1, step=0.3) # Fixed step size
      np.arange(12).reshape(3,4) # Create and reshape
      np.eye(4) # Identity matrix
```

```
array([ 0.    ,  0.333,  0.667,  1.    ])
```

```
array([ 0. ,  0.3,  0.6,  0.9])
```

```
array([[ 0,  1,  2,  3],  
       [ 4,  5,  6,  7],  
       [ 8,  9, 10, 11]])
```

```
array([[ 1.,  0.,  0.,  0.],  
       [ 0.,  1.,  0.,  0.],  
       [ 0.,  0.,  1.,  0.],  
       [ 0.,  0.,  0.,  1.]])
```

Basic Operations

Arithmetic operators on arrays apply elementwise. A new array is created and filled with the result. Some operations, such as += and *=, act in place to modify an existing array rather than create a new one.

```
[8]: a = np.array([20, 30, 40, 50])  
     b = np.arange(4)  
     a, b      # Just printing  
     a-b  
     b**2  
     a > 32  
     a += 1  
     a
```

```
(array([20, 30, 40, 50]), array([0, 1, 2, 3]))
```

```
array([20, 29, 38, 47])
```

```
array([0, 1, 4, 9])
```

```
array([False, False,  True,  True], dtype=bool)
```

```
array([21, 31, 41, 51])
```

The product operator * operates elementwise.
The matrix product can be performed using dot()

```
[9]: A, B = np.array([[1,1], [0,1]]), np.array([[2,0], [3,4]]) # assign multiple  
     A  
     B  
     A * B  
     np.dot(A, B)
```

```
array([[1, 1],
       [0, 1]])
```

```
array([[2, 0],
       [3, 4]])
```

```
array([[2, 0],
       [0, 4]])
```

```
array([[5, 4],
       [3, 4]])
```

Upcasting: Operations with arrays of different types choose the more general/precise one.

```
[10]: a = np.ones(3, dtype=np.int) # initialize to integers
      b = np.linspace(0, np.pi, 3) # default type is float
      a.dtype, b.dtype, (a + b).dtype

(dtype('int64'), dtype('float64'), dtype('float64'))
```

ndarrays have most unary operations (max,min,sum,...) built in

```
[11]: a = np.random.random((2,3))
      a
      a.sum(), a.min(), a.max()

array([[ 0.922,  0.843,  0.676],
       [ 0.22 ,  0.97 ,  0.599]])

(4.2293775981866508, 0.22019608034032734, 0.96968406069707935)
```

By specifying the axis parameter you can apply an operation along a specified axis of an array

```
[12]: b = np.arange(12).reshape(3,4)
      b
      b.sum(axis=0)
      b.sum(axis=1)

array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])

array([12, 15, 18, 21])

array([ 6, 22, 38])
```

Universal Functions

NumPy provides familiar mathematical functions such as `sin`, `cos`, `exp`, `sqrt`, `floor`,... In NumPy, these are called “universal functions” (ufunc), and operate elementwise on an array, producing an array as output.

```
[13]: np.sin(np.arange(0, 10))  
  
array([ 0.      ,  0.841,  0.909,  0.141, -0.757, -0.959, -0.279,  0.657,  
        0.989,  0.412])
```

Shape Manipulation

Transpose, flatten, reshape,...

```
[14]: a = np.floor(10*np.random.random((3,4)))  
      a  
      a.transpose()  
      b = a.ravel() # flatten array  
      b  
      b.reshape(3, -1) # reshape in 2 rows (and as many columns as needed)  
  
array([[ 6.,  5.,  2.,  5.],  
       [ 1.,  8.,  2.,  3.],  
       [ 2.,  6.,  3.,  8.]])  
  
array([[ 6.,  1.,  2.],  
       [ 5.,  8.,  6.],  
       [ 2.,  2.,  3.],  
       [ 5.,  3.,  8.]])  
  
array([ 6.,  5.,  2.,  5.,  1.,  8.,  2.,  3.,  2.,  6.,  3.,  8.])  
  
array([[ 6.,  5.,  2.,  5.],  
       [ 1.,  8.,  2.,  3.],  
       [ 2.,  6.,  3.,  8.]])
```

Arrays can be split and stacked together

```
[15]: a = np.floor(10*np.random.random((2,6)))  
      a  
      b, c = np.hsplit(a, 2) # Idem: vsplit for vertical splits  
      b  
      c  
      np.hstack((b, c)) # Idem: vstack for vertical stacks  
  
array([[ 3.,  4.,  4.,  6.,  8.,  3.],  
       [ 2.,  9.,  6.,  5.,  3.,  1.]])
```

```
array([[ 3.,  4.,  4.],
       [ 2.,  9.,  6.]])
```

```
array([[ 6.,  8.,  3.],
       [ 5.,  3.,  1.]])
```

```
array([[ 3.,  4.,  4.,  6.,  8.,  3.],
       [ 2.,  9.,  6.,  5.,  3.,  1.]])
```

Indexing and Slicing

Arrays can be indexed and sliced using [start:stop:stepsize]. Defaults are [0:ndim:1]

```
[16]: a = np.arange(10)**2
      a
```

```
array([ 0,  1,  4,  9, 16, 25, 36, 49, 64, 81])
```

```
[17]: a[2]
```

```
4
```

```
[18]: a[3:10:2]
```

```
array([ 9, 25, 49, 81])
```

```
[19]: a[::-1] # Defaults are used if indices not stated
```

```
array([81, 64, 49, 36, 25, 16,  9,  4,  1,  0])
```

```
[20]: a[::2]
```

```
array([ 0,  4, 16, 36, 64])
```

For multi-dimensional arrays, axes are comma-separated: [x,y,z].

```
[21]: b = np.arange(16).reshape(4,4)
      b
      b[2,3] # row 2, column 3
```

```
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11],
       [12, 13, 14, 15]])
```

11

```
[22]: b[0:3,1] # Values 0 to 3 in column 1
      b[ : ,1] # The whole column 1
```

```
array([1, 5, 9])
```

```
array([ 1,  5,  9, 13])
```

```
[23]: b[1:3, : ] # Rows 1:3, all columns
```

```
array([[ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
```

```
[24]: # Return the last row
      b[-1]
```

```
array([12, 13, 14, 15])
```

Note: dots (...) represent as many colons (:) as needed * $x[1,2,\dots] = x[1,2,:::] * x[\dots,3] = x[:,:::,3] * x[4,\dots,5,:] = x[4,::,5,:]$

Arrays can also be indexed by arrays of integers and booleans.

```
[25]: a = np.arange(12)**2
      i = np.array([ 1,1,3,8,5 ])
      a
      a[i]
```

```
array([ 0,  1,  4,  9, 16, 25, 36, 49, 64, 81, 100, 121])
```

```
array([ 1,  1,  9, 64, 25])
```

A matrix of indices returns a matrix with the corresponding values.

```
[26]: j = np.array([[ 3, 4], [9, 7]])
      a[j]
```

```
array([[ 9, 16],
       [81, 49]])
```

With boolean indices we explicitly choose which items in the array we want and which ones we don't.

```
[27]: a = np.arange(12).reshape(3,4)
      a
      a[np.array([False, True, True]), :]
      b = a > 4
      b
      a[b]
```

```
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
```

```
array([[ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
```

```
array([[False, False, False, False],
       [False,  True,  True,  True],
       [ True,  True,  True,  True]], dtype=bool)
```

```
array([ 5,  6,  7,  8,  9, 10, 11])
```

Iterating

Iterating is done with respect to the first axis:

```
[28]: for row in b:
        print(row)

[False False False False]
[False  True  True  True]
[ True  True  True  True]
```

Operations on each element can be done by flattening the array (or nested loops)

```
[29]: for element in b.flat: # flat returns an iterator
        print(element)
```

```
False
False
False
False
False
True
True
True
True
True
True
True
True
```

Copies and Views (or: how to shoot yourself in a foot)

Assigning an array to another variable does NOT create a copy

```
[30]: a = np.arange(12)
        b = a
        a
```



```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
[31]: b[0] = -100
      b
```

```
array([-100,    1,    2,    3,    4,    5,    6,    7,    8,    9,   10,
        11])
```

```
[32]: a
```

```
array([-100,    1,    2,    3,    4,    5,    6,    7,    8,    9,   10,
        11])
```

The `view()` method creates a NEW array object that looks at the same data.

```
[33]: a = np.arange(12)
      a
      c = a.view()
      c.resize((2, 6))
      c
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
array([[ 0,  1,  2,  3,  4,  5],
       [ 6,  7,  8,  9, 10, 11]])
```

```
[34]: a[0] = 123
      c # c is also changed now
```

```
array([[123,    1,    2,    3,    4,    5],
       [ 6,    7,    8,    9,   10,   11]])
```

Slicing an array returns a view of it.

```
[35]: c
      s = c[ : , 1:3]
      s[:] = 10
      s
      c
```

```
array([[123,    1,    2,    3,    4,    5],
       [ 6,    7,    8,    9,   10,   11]])
```

```
array([[10, 10],
       [10, 10]])
```

```
array([[123, 10, 10, 3, 4, 5],
       [ 6, 10, 10, 9, 10, 11]])
```

The `copy()` method makes a deep copy of the array and its data.

```
[36]: d = a.copy()
      d[0] = -42
      d
array([-42, 10, 10, 3, 4, 5, 6, 10, 10, 9, 10, 11])
```

```
[37]: a
array([123, 10, 10, 3, 4, 5, 6, 10, 10, 9, 10, 11])
```

Numpy: further reading

- Numpy Tutorial: http://wiki.scipy.org/Tentative_NumPy_Tutorial
- “Python for Data Analysis” by Wes McKinney (O’Reilly)

SciPy

SciPy is a collection of packages for scientific computing, among others:

- `scipy.integrate`: numerical integration and differential equation solvers
- `scipy.linalg`: linear algebra routines and matrix decompositions
- `scipy.optimize`: function optimizers (minimizers) and root finding algorithms
- `scipy.signal`: signal processing tools
- `scipy.sparse`: sparse matrices and sparse linear system solvers
- `scipy.stats`: probability distributions, statistical tests, descriptive statistics

Sparse matrices

Sparse matrices are used in scikit-learn for (large) arrays that contain mostly zeros. You can convert a dense (numpy) matrix to a sparse matrix.

```
[38]: from scipy import sparse
      eye = np.eye(4)
      eye
      sparse_matrix = sparse.csr_matrix(eye) # Compressed Sparse Row matrix
      sparse_matrix
      print("{}".format(sparse_matrix))
array([[ 1.,  0.,  0.,  0.],
       [ 0.,  1.,  0.,  0.],
       [ 0.,  0.,  1.,  0.],
       [ 0.,  0.,  0.,  1.]])
```

```
<4x4 sparse matrix of type '<class 'numpy.float64'>'
with 4 stored elements in Compressed Sparse Row format>
```

```
(0, 0)      1.0
(1, 1)      1.0
(2, 2)      1.0
(3, 3)      1.0
```

When the data is too large, you can create a sparse matrix by passing the values and coordinates (COO format).

```
[39]: data = np.ones(4) # [1,1,1,1]
      row_indices = col_indices = np.arange(4) # [1,2,3,4]
      eye_coo = sparse.coo_matrix((data, (row_indices, col_indices)))
      print("{}".format(eye_coo))
```

```
(0, 0)      1.0
(1, 1)      1.0
(2, 2)      1.0
(3, 3)      1.0
```

Further reading

Check the [SciPy reference guide](#) for tutorials and examples of all SciPy capabilities.

pandas

pandas is a Python library for data wrangling and analysis. It provides:

- DataFrame: a table, similar to an R DataFrame that holds any structured data
 - Every column can have its own data type (strings, dates, floats,...)
- A great range of methods to apply to this table (sorting, querying, joining,...)
- Imports data from a wide range of data formats (CVS, Excel) and databases (e.g. SQL)

Series

A one-dimensional array of data (of any numpy type), with indexed values. It can be created by passing a Python list or dict, a numpy array, a csv file,...

```
[40]: import pandas as pd
      pd.Series([1,3,np.nan]) # Default integers are integers
      pd.Series([1,3,5], index=['a','b','c'])
      pd.Series({'a': 1, 'b': 2, 'c': 3 }) # when given a dict, the keys will b
      pd.Series({'a': 1, 'b': 2, 'c': 3 }, index = ['b', 'c', 'd']) # this will
```

```
0    1.0
1    3.0
2    NaN
dtype: float64
```

```
a    1
b    3
c    5
dtype: int64
```

```
a    1
b    2
c    3
dtype: int64
```

```
b    2.0
c    3.0
d    NaN
dtype: float64
```

Functions like a numpy array, however with index labels as indices

```
[41]: a = pd.Series({'a' : 1, 'b': 2, 'c': 3 })
      a
      a['b']          # Retrieves a value
      a[['a','b']]    # and can also be sliced
```

```
a    1
b    2
c    3
dtype: int64
```

```
2
```

```
a    1
b    2
dtype: int64
```

numpy array operations on Series preserve the index value

```
[42]: a
      a[a > 1]
      a * 2
      np.sqrt(a)
```

```
a    1
b    2
c    3
dtype: int64
```

```
b      2
c      3
dtype: int64
```

```
a      2
b      4
c      6
dtype: int64
```

```
a      1.00
b      1.41
c      1.73
dtype: float64
```

Operations over multiple Series will align the indices

```
[43]: a = pd.Series({'John' : 1000, 'Mary': 2000, 'Andre': 3000 })
      b = pd.Series({'John' : 100, 'Andre': 200, 'Cecilia': 300 })
      a + b
```

```
Andre      3200.0
Cecilia      NaN
John       1100.0
Mary        NaN
dtype: float64
```

DataFrame

A DataFrame is a tabular data structure with both a row and a column index. It can be created by passing a dict of arrays, a csv file,...

```
[44]: data = {'state': ['Ohio', 'Ohio', 'Nevada', 'Nevada'], 'year': [2000, 2001, 2001, 2002],
              'pop': [1.5, 1.7, 2.4, 2.9]}
      pd.DataFrame(data)
      pd.DataFrame(data, columns=['year', 'state', 'pop', 'color']) # Will match
```

```
   pop  state  year
0  1.5   Ohio  2000
1  1.7   Ohio  2001
2  2.4  Nevada  2001
3  2.9  Nevada  2002
```

```
   year  state  pop  color
0  2000   Ohio  1.5   NaN
1  2001   Ohio  1.7   NaN
2  2001  Nevada  2.4   NaN
3  2002  Nevada  2.9   NaN
```

It can be composed with a numpy array and row and column indices, and decomposed

```
[45]: dates = pd.date_range('20130101', periods=4)
      df = pd.DataFrame(np.random.randn(4, 4), index=dates, columns=list('ABCD'))
      df
```

```
          A      B      C      D
2013-01-01 -2.42  1.89  0.53 -0.97
2013-01-02 -0.79  1.12  0.42 -0.35
2013-01-03 -0.38 -2.38 -0.94  0.39
2013-01-04  0.28 -1.16  0.09  1.54
```

```
[46]: df.index
      df.columns
      df.values
```

```
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04'], dtype='datetime64[ns]', freq='D')
```

```
Index(['A', 'B', 'C', 'D'], dtype='object')
```

```
array([[ -2.417,  1.887,  0.526, -0.967],
       [ -0.794,  1.121,  0.422, -0.35 ],
       [ -0.384, -2.378, -0.935,  0.393],
       [  0.275, -1.162,  0.088,  1.543]])
```

DataFrames can easily read/write data from/to files

- `read_csv(source)`: load CSV data from file or url
- `read_table(source, sep=',')`: load delimited data with separator
- `df.to_csv(target)`: writes the DataFrame to a file

```
[47]: dfs = pd.read_csv('data.csv')
      dfs
      dfs.set_value(0, 'a', 10)
      dfs.to_csv('data.csv', index=False) # Don't export the row index
```

```
   a    b    c    d message
0  10    2    3    4   hello
1   5    6    7    8   world
2   9   10   11   12     foo
```

```
   a    b    c    d message
0  10    2    3    4   hello
1   5    6    7    8   world
2   9   10   11   12     foo
```

Simple operations

```
[48]: df.head() # First 5 rows
      df.tail() # Last 5 rows
```

	A	B	C	D
2013-01-01	-2.42	1.89	0.53	-0.97
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54

	A	B	C	D
2013-01-01	-2.42	1.89	0.53	-0.97
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54

```
[49]: # Quick stats
      df.describe()
```

	A	B	C	D
count	4.00	4.00	4.00	4.00
mean	-0.83	-0.13	0.03	0.15
std	1.15	1.98	0.67	1.08
min	-2.42	-2.38	-0.94	-0.97
25%	-1.20	-1.47	-0.17	-0.50
50%	-0.59	-0.02	0.25	0.02
75%	-0.22	1.31	0.45	0.68
max	0.28	1.89	0.53	1.54

```
[50]: # Transpose
      df.T
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04
A	-2.42	-0.79	-0.38	0.28
B	1.89	1.12	-2.38	-1.16
C	0.53	0.42	-0.94	0.09
D	-0.97	-0.35	0.39	1.54

```
[51]: df.sort_index(axis=1, ascending=False) # Sort by index labels
      df.sort(columns='B') # Sort by values
```

	D	C	B	A
2013-01-01	-0.97	0.53	1.89	-2.42
2013-01-02	-0.35	0.42	1.12	-0.79
2013-01-03	0.39	-0.94	-2.38	-0.38
2013-01-04	1.54	0.09	-1.16	0.28

	A	B	C	D
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-01	-2.42	1.89	0.53	-0.97

Selecting and slicing

```
[52]: df['A'] # Get single column by label
      df.A    # Shorthand
```

2013-01-01	-2.42
2013-01-02	-0.79
2013-01-03	-0.38
2013-01-04	0.28

Freq: D, Name: A, dtype: float64

2013-01-01	-2.42
2013-01-02	-0.79
2013-01-03	-0.38
2013-01-04	0.28

Freq: D, Name: A, dtype: float64

```
[53]: df[0:2] # Get rows by index number
      df.iloc[0:2,0:2] # Get rows and columns by index number
      df['20130102':'20130103'] # or row label
      df.loc['20130102':'20130103', ['A','B']] # or row and column label
      df.ix[0:2, ['A','B']] # allows mixing integers and labels
```

	A	B	C	D
2013-01-01	-2.42	1.89	0.53	-0.97
2013-01-02	-0.79	1.12	0.42	-0.35

	A	B
2013-01-01	-2.42	1.89
2013-01-02	-0.79	1.12

	A	B	C	D
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-03	-0.38	-2.38	-0.94	0.39

	A	B
2013-01-02	-0.79	1.12
2013-01-03	-0.38	-2.38

	A	B
2013-01-01	-2.42	1.89
2013-01-02	-0.79	1.12

`query()` retrieves data matching a boolean expression

```
[54]: df
      df.query('A > 0.4') # Identical to df[df.A > 0.4]
      df.query('A > B')   # Identical to df[df.A > df.B]
```

	A	B	C	D
2013-01-01	-2.42	1.89	0.53	-0.97
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54

```
Empty DataFrame
Columns: [A, B, C, D]
Index: []
```

	A	B	C	D
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54

Note: similar to NumPy, indexing and slicing returns a *view* on the data. Use `copy()` to make a deep copy.

Operations

DataFrames offer a [wide range of operations](#): max, mean, min, sum, std,...

```
[55]: df.mean()           # Mean of all values per column
      df.mean(axis=1)     # Other axis: means per row
```

```
A    -0.83
B    -0.13
C     0.03
D     0.15
dtype: float64
```

```
2013-01-01    -0.24
2013-01-02     0.10
2013-01-03    -0.83
2013-01-04     0.19
Freq: D, dtype: float64
```

All of numpy's universal functions also work with dataframes

```
[56]: np.abs(df)
```

	A	B	C	D
2013-01-01	2.42	1.89	0.53	0.97
2013-01-02	0.79	1.12	0.42	0.35
2013-01-03	0.38	2.38	0.94	0.39
2013-01-04	0.28	1.16	0.09	1.54

Other (custom) functions can be applied with `apply(func)`

```
[57]: df
      df.apply(np.max)
      df.apply(lambda x: x.max() - x.min())
```

	A	B	C	D
2013-01-01	-2.42	1.89	0.53	-0.97
2013-01-02	-0.79	1.12	0.42	-0.35
2013-01-03	-0.38	-2.38	-0.94	0.39
2013-01-04	0.28	-1.16	0.09	1.54

```
A      0.28
B      1.89
C      0.53
D      1.54
dtype: float64
```

```
A      2.69
B      4.26
C      1.46
D      2.51
dtype: float64
```

Data can be aggregated with `groupby()`

```
[58]: df = pd.DataFrame({'A' : ['foo', 'bar', 'foo', 'bar'], 'B' : ['one', 'one',
                                'C' : np.random.randn(4), 'D' : np.random.randn(4) })
      df
      df.groupby('A').sum()
      df.groupby(['A', 'B']).sum()
```

	A	B	C	D
0	foo	one	0.34	1.30
1	bar	one	0.84	-0.05
2	foo	two	-0.95	-1.67
3	bar	three	-0.33	1.45

	C	D
A		
bar	0.50	1.40
foo	-0.61	-0.38

		C	D
A	B		
bar	one	0.84	-0.05
	three	-0.33	1.45
foo	one	0.34	1.30
	two	-0.95	-1.67

Data wrangling (some examples)

Merge: combine two dataframes based on common keys

```
[59]: df1 = pd.DataFrame({'key': ['b', 'b', 'a'], 'data1': range(3)})
      df2 = pd.DataFrame({'key': ['a', 'b'], 'data2': range(2)})
      df1
      df2
      pd.merge(df1, df2)
```

	data1	key
0	0	b
1	1	b
2	2	a

	data2	key
0	0	a
1	1	b

	data1	key	data2
0	0	b	1
1	1	b	1
2	2	a	0

Append: append one dataframe to another

```
[60]: df = pd.DataFrame(np.random.randn(2, 4))
      df
      s = pd.DataFrame(np.random.randn(1, 4))
      s
      df.append(s, ignore_index=True)
```

	0	1	2	3
0	-0.17	-0.81	1.95	-1.08
1	-0.32	1.73	-0.48	0.90

	0	1	2	3
0	-0.34	0.32	1.16	-1.03

```

      0      1      2      3
0 -0.17 -0.81  1.95 -1.08
1 -0.32  1.73 -0.48  0.90
2 -0.34  0.32  1.16 -1.03

```

Remove duplicates

```

[61]: df = pd.DataFrame({'k1': ['one'] * 3, 'k2': [1, 1, 2]})
      df
      df.drop_duplicates()

```

```

      k1  k2
0  one   1
1  one   1
2  one   2

```

```

      k1  k2
0  one   1
2  one   2

```

Replace values

```

[62]: df = pd.DataFrame({'k1': [1, -1], 'k2': [-1, 2]}) # Say that -1 is a senti
      df
      df.replace(-1, np.nan)

```

```

      k1  k2
0     1  -1
1    -1   2

```

```

      k1  k2
0  1.0 NaN
1  NaN  2.0

```

Discretization and binning

```

[63]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
      bins = [18, 25, 35, 60, 100]
      cats = pd.cut(ages, bins)
      cats.labels
      pd.value_counts(cats)

```

```

array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)

```

```

(18, 25]      5
(35, 60]      3
(25, 35]      3
(60, 100]     1
dtype: int64

```

Further reading

- Pandas docs: <http://pandas.pydata.org/pandas-docs/stable/>
- <https://bitbucket.org/hrojas/learn-pandas>
- Python for Data Analysis (O'Reilly) by Wes McKinney (the author of pandas)

matplotlib

[matplotlib](#) is the primary scientific plotting library in Python. It provides:

- Publication-quality [visualizations](#) such as line charts, histograms, and scatter plots.
- Integration in pandas to make plotting much easier.
- Interactive plotting in Jupyter notebooks for quick visualizations.
 - Requires some setup. See preamble and [%matplotlib](#).
- Many GUI backends, export to PDF, SVG, JPG, PNG, BMP, GIF, etc.
- Ecosystem of libraries for more advanced plotting, e.g. [Seaborn](#)


Low-level usage

`plot()` is the [main function](#) to generate a plot (but many more exist):

```
plot(x, y)           Plot x vs y, default settings
plot(x, y, 'bo')     Plot x vs y, blue circle markers
plot(y, 'r+')        Plot y (x = array 0..N-1), red plusses
```

Every plotting function is completely customizable through a large set of options.


```
[71]: x = np.linspace(-10, 10, 100) # Sequence for X-axis
      y = np.sin(x) # sine values
      p = plt.plot(x, y, marker="x") # Line plot with marker x
```




pandas + matplotlib

pandas DataFrames offer an easier, higher-level interface for matplotlib functions


```
[65]: df = pd.DataFrame(np.random.randn(500, 4),  
                        columns=['a', 'b', 'c', 'd']) # random 4D data  
p = df.cumsum().plot() # Plot cumulative sum of all series
```



```
[66]: p = df[:10].plot(kind='bar') # First 10 arrays as bar plots
```




```
[67]: p = df.boxplot() # Boxplot for each of the 4 series
```



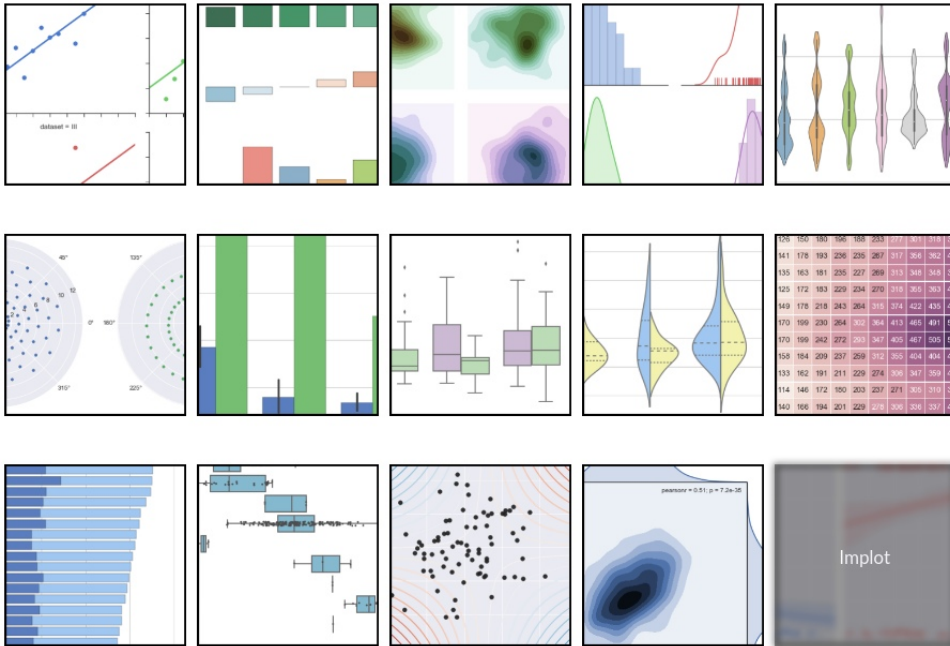
```
[68]: # Scatter plot using the 4 series for x, y, color, scale
df[:300].plot(kind='scatter', x='a', y='b', c='c',
                s=df['d']*50, linewidth='0', cmap='plasma')
```

<matplotlib.axes._subplots.AxesSubplot at 0x111f72400>



Advanced plotting libraries

Several libraries, such as [Seaborn](#) offer more advanced plots and easier interfaces.



Further reading links

- [Matplotlib examples](#)
- [Plotting with pandas](#)
- [Seaborn examples](#)