

OpenML in Python

OpenML is an online collaboration platform for machine learning:

- Find or share interesting, well-documented datasets
- Define research / modelling goals (tasks)
- Explore large amounts of machine learning algorithms, with APIs in Java, R, Python
- Log and share reproducible experiments, models, results
- Works seamlessly with scikit-learn and other libraries
- Large scale benchmarking, compare to state of the art

Installation

Install the OpenML developer version

```
'pip install openml' coming up (october 2017) pip install  
git+https://github.com/renatopp/liac-arff@master  
git+https://github.com/openml/openml-python.git@develop
```

```
[1]: from IPython.display import set_matplotlib_formats, display, HTML  
HTML(''<style>html, body{overflow-y: visible !important} .CodeMirror{min-w  
<IPython.core.display.HTML object>
```

Authentication

- Create an OpenML account (free) on <http://www.openml.org>.
- After logging in, open your account page (avatar on the top right)
- Open 'Account Settings', then 'API authentication' to find your API key.

There are two ways to authenticate:

- Create a plain text file `~/.openml/config` with the line `'apikey=MYKEY'`, replacing MYKEY with your API key.
- Run the code below, replacing `'YOURKEY'` with your API key.

```
[2]: # Uncomment and set your OpenML key. Don't share your key with others.  
import openml as oml  
#oml.config.apikey = 'YOURKEY'
```

It all starts with data

Explore thousands of datasets, or share your own

List datasets

```
[3]: import openml as oml  
openml_list = oml.datasets.list_datasets() # Returns a dict  
  
# Show a nice table with some key data properties  
import pandas as pd  
datalist = pd.DataFrame.from_dict(openml_list, orient='index')  
datalist = datalist[['did', 'name', 'NumberOfInstances',
```

```

        'NumberOfFeatures', 'NumberOfClasses']]
print("First 10 of %s datasets..." % len(datalist))
datalist[:10]

```

First 10 of 19530 datasets...

	did	name	NumberOfInstances	NumberOfFeatures	NumberOfClasses
1	1	anneal	898.0	39.0	6.0
2	2	anneal	898.0	39.0	5.0
3	3	kr-vs-kp	3196.0	37.0	2.0
4	4	labor	57.0	17.0	2.0
5	5	arrhythmia	452.0	280.0	13.0
6	6	letter	20000.0	17.0	26.0
7	7	audiology	226.0	70.0	24.0
8	8	liver-disorders	345.0	7.0	-1.0
9	9	autos	205.0	26.0	6.0
10	10	lymph	148.0	19.0	4.0

Exercise

- Find datasets with more than 10000 examples
- Find a dataset called 'eeg_eye_state'
- Find all datasets with more than 50 classes

```

[4]: datalist[datalist.NumberOfInstances>10000
      ].sort(['NumberOfInstances'])[:20]

```

	did	name	NumberOfInstances	\
23515	23515	sulfur	10081.0	
372	372	internet_usage	10108.0	
981	981	kdd_internet_usage	10108.0	
1536	1536	volcanoes-b6	10130.0	
4562	4562	InternetUsage	10168.0	
1531	1531	volcanoes-b1	10176.0	
1534	1534	volcanoes-b4	10190.0	
1459	1459	artificial-characters	10218.0	
1478	1478	har	10299.0	
1533	1533	volcanoes-b3	10386.0	
1532	1532	volcanoes-b2	10668.0	
1053	1053	jm1	10885.0	
1414	1414	Kaggle_bike_sharing_demand_challenge	10886.0	
1044	1044	eye_movements	10936.0	
32	32	pendigits	10992.0	
1019	1019	pendigits	10992.0	
4534	4534	PhishingWebsites	11055.0	
399	399	ohscal.wc	11162.0	
310	310	mammography	11183.0	
1568	1568	nursery	12958.0	

NumberOfFeatures NumberOfClasses

23515	7.0	-1.0
372	72.0	46.0
981	69.0	2.0
1536	4.0	5.0
4562	72.0	-1.0
1531	4.0	5.0
1534	4.0	5.0
1459	8.0	10.0
1478	562.0	6.0
1533	4.0	5.0
1532	4.0	5.0
1053	22.0	2.0
1414	12.0	-1.0
1044	28.0	3.0
32	17.0	10.0
1019	17.0	2.0
4534	31.0	2.0
399	11466.0	10.0
310	7.0	2.0
1568	9.0	4.0

```
[5]: datalist.query('name == "MagicTelescope"')
```

	did	name	NumberOfInstances	NumberOfFeatures	\
1120	1120	MagicTelescope	19020.0	12.0	
		NumberOfClasses			
1120		2.0			

```
[6]: datalist.query('NumberOfClasses > 50')
```

	did	name	NumberOfInstances	NumberOfFeatures	\
1491	1491	one-hundred-plants-margin	1600.0	65.0	
1492	1492	one-hundred-plants-shape	1600.0	65.0	
1493	1493	one-hundred-plants-texture	1599.0	65.0	
4546	4546	Plants	44940.0	16.0	
4552	4552	BachChoralHarmony	5665.0	17.0	
40753	40753	delays_zurich_transport	5465575.0	15.0	
		NumberOfClasses			
1491		100.0			
1492		100.0			
1493		100.0			
4546		57.0			
4552		102.0			
40753		4082.0			

Download datasets

This is done based on the dataset ID ('did').

```
[7]: dataset = oml.datasets.get_dataset(1120)

# Print a summary
print("This is dataset '%s', the target feature is '%s'" %
      (dataset.name, dataset.default_target_attribute))
print("URL: %s" % dataset.url)
print(dataset.description[:500])
```

This is dataset 'MagicTelescope', the target feature is 'class:'

URL: <https://www.openml.org/data/v1/download/54003/MagicTelescope.ARF>

****Author**:** R. K. Bock. Major Atmospheric Gamma Imaging Cherenkov Telescope proj

Donated by P. Savicky, Institute of Computer Science, AS of CR, Czech Republic

****Source**:** [UCI] (<https://archive.ics.uci.edu/ml/datasets/magic+gamma+telescope>)

****Please cite**:**

The data are MC generated (see below) to simulate registration of high energy ga

Get the actual data.

Returned as numpy array, with meta-info (e.g. target feature, feature names,...)

```
[8]: X, y, attribute_names = dataset.get_data(
      target=dataset.default_target_attribute,
      return_attribute_names=True)
eeg = pd.DataFrame(X, columns=attribute_names)
eeg['class'] = y
print(eeg[:10])
```

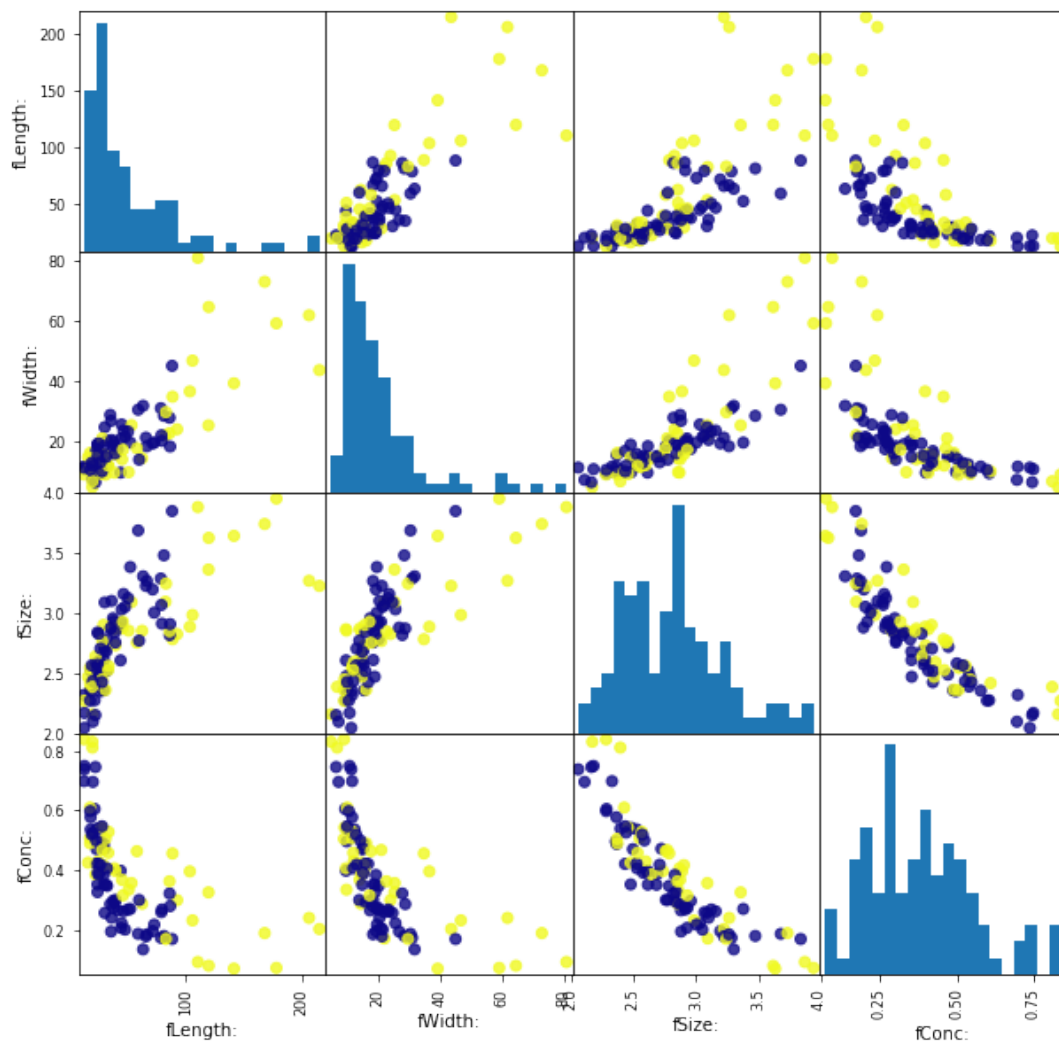
	fLength:	fWidth:	fSize:	fConc:	fConc1:	fAsym:	fM3Long:	\
0	28.796700	16.002100	2.6449	0.3918	0.1982	27.700399	22.011000	
1	31.603600	11.723500	2.5185	0.5303	0.3773	26.272200	23.823799	
2	162.052002	136.031006	4.0612	0.0374	0.0187	116.740997	-64.858002	
3	23.817200	9.572800	2.3385	0.6147	0.3922	27.210699	-6.463300	
4	75.136200	30.920500	3.1611	0.3168	0.1832	-5.527700	28.552500	
5	51.624001	21.150200	2.9085	0.2420	0.1340	50.876099	43.188702	
6	48.246799	17.356501	3.0332	0.2529	0.1515	8.573000	38.095699	
7	26.789700	13.759500	2.5521	0.4236	0.2174	29.633900	20.455999	
8	96.232697	46.516499	4.1540	0.0779	0.0390	110.355003	85.048599	
9	46.761902	15.199300	2.5786	0.3377	0.1913	24.754801	43.877102	

	fM3Trans:	fAlpha:	fDist:	class
0	-8.202700	40.091999	81.882797	0
1	-9.957400	6.360900	205.261002	0
2	-45.216000	76.959999	256.787994	0
3	-7.151300	10.449000	116.737000	0
4	21.839300	4.648000	356.462006	0
5	9.814500	3.613000	238.098007	0
6	10.586800	4.792000	219.087006	0
7	-2.929200	0.812000	237.134003	0
8	43.184399	4.854000	248.225998	0
9	-6.681200	7.875000	102.250999	0

Exercise

- Explore the data visually

```
[9]: %matplotlib inline
eegs = eeg.sample(n=1000)
pd.scatter_matrix(eegs.iloc[:100,:4], c=eegs[:100]['class'], figsize=(10, 10),
                  marker='o', hist_kwds={'bins': 20},
                  alpha=.8, cmap='plasma');
```



Train machine learning models

Train a scikit-learn model on the data manually

```
[10]: from sklearn import neighbors

dataset = oml.datasets.get_dataset(1120)
X, y = dataset.get_data(target=dataset.default_target_attribute)
clf = neighbors.KNeighborsClassifier(n_neighbors=1)
clf.fit(X, y)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                     weights='uniform')
```

You can also ask for meta-data to automatically preprocess the data - e.g. categorical features
-> do feature encoding

```
[11]: from sklearn import preprocessing
      dataset = oml.datasets.get_dataset(10)
      X, y, categorical = dataset.get_data(
          target=dataset.default_target_attribute,
          return_categorical_indicator=True)
      print("Categorical features: %s" % categorical)
      enc = preprocessing.OneHotEncoder(categorical_features=categorical)
      X = enc.fit_transform(X)
      clf.fit(X, y)
```

```
Categorical features: [True, True, True, True, True, True, True, True, True, False, Fa
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                     weights='uniform')
```

Tasks: set your own goals

and invite others to work on the same problem

Note: tasks are typically created in the web interface

Listing tasks

```
[12]: task_list = oml.tasks.list_tasks(size=5000) # Get first 5000 tasks

      mytasks = pd.DataFrame(task_list).transpose()
      mytasks = mytasks[['tid', 'did', 'name', 'task_type', 'estimation_procedure', '
      #print(mytasks.columns)
      print("First 5 of %s tasks:" % len(mytasks))
      print(mytasks.head())
```

First 5 of 5000 tasks:

	tid	did	name	task_type	estimation_procedure	\
1	1	1	anneal	Supervised Classification	10-fold Crossvalidation	
2	2	2	anneal	Supervised Classification	10-fold Crossvalidation	
3	3	3	kr-vs-kp	Supervised Classification	10-fold Crossvalidation	
4	4	4	labor	Supervised Classification	10-fold Crossvalidation	
5	5	5	arrhythmia	Supervised Classification	10-fold Crossvalidation	

```
      evaluation_measures
1  predictive_accuracy
2  predictive_accuracy
```

```
3 predictive_accuracy
4 predictive_accuracy
5 predictive_accuracy
```

Exercise

Search for the tasks on the 'eeg-eye-state' dataset

```
[13]: print(mytasks.query('name=="MagicTelescope"))
```

	tid	did	name	task_type	\
3954	3954	1120	MagicTelescope	Supervised Classification	
4659	4659	1120	MagicTelescope	Supervised Classification	
7228	7228	1120	MagicTelescope	Supervised Data Stream Classification	
10067	10067	1120	MagicTelescope	Learning Curve	

	estimation_procedure	evaluation_measures
3954	10-fold Crossvalidation	predictive_accuracy
4659	10 times 10-fold Crossvalidation	predictive_accuracy
7228	Interleaved Test then Train	NaN
10067	10-fold Learning Curve	NaN

Download tasks

```
[14]: from pprint import pprint
task = oml.tasks.get_task(3954)
pprint(vars(task))
```

```
{'class_labels': ['g', 'h'],
 'cost_matrix': None,
 'dataset_id': 1120,
 'estimation_parameters': {'number_folds': '10',
                           'number_repeats': '1',
                           'percentage': '',
                           'stratified_sampling': 'true'},
 'estimation_procedure': {'data_splits_url': 'https://www.openml.org/api_splits/
                           'parameters': {'number_folds': '10',
                                           'number_repeats': '1',
                                           'percentage': '',
                                           'stratified_sampling': 'true'},
                           'type': 'crossvalidation'},
 'evaluation_measure': 'predictive_accuracy',
 'split': None,
 'target_name': 'class:',
 'task_id': 3954,
 'task_type': 'Supervised Classification',
 'task_type_id': 1}
```

Runs: Easily explore models by running them on tasks

We can run (many) scikit-learn algorithms on (many) OpenML tasks.

```
[15]: from sklearn import ensemble

      # Get a task
      task = oml.tasks.get_task(3954)

      # Build any classifier or pipeline
      clf = ensemble.RandomForestClassifier()

      # Create a flow
      flow = oml.flows.sklearn_to_flow(clf)

      # Run the flow
      run = oml.runs.run_flow_on_task(task, flow)
```

Share the run on the OpenML server

```
[16]: myrun = run.publish()
      print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))
```

Uploaded to <http://www.openml.org/r/6068436>

It also works with pipelines

When you need to handle ‘dirty’ data, build pipelines to model then automatically

```
[17]: from sklearn import pipeline, ensemble, preprocessing
      from openml import tasks, runs, datasets
      task = tasks.get_task(59)
      pipe = pipeline.Pipeline(steps=[
          ('Imputer', preprocessing.Imputer(strategy='median')),
          ('OneHotEncoder', preprocessing.OneHotEncoder(sparse=False, ha
          ('Classifier', ensemble.RandomForestClassifier())
      ])
      flow = oml.flows.sklearn_to_flow(pipe)

      run = oml.runs.run_flow_on_task(task, flow)
      myrun = run.publish()
      print("Uploaded to http://www.openml.org/r/" + str(myrun.run_id))
```

Uploaded to <http://www.openml.org/r/6068437>

Download previous results

You can download all your results anytime, as well as everybody else’s

```
[40]: import seaborn as sns
      import pandas as pd
```



```


from matplotlib import pyplot
from openml import evaluations

# Get the list of runs for task 3954
evaluations = oml.evaluations.list_evaluations(task=[3954], function='area

# Download the tasks and plot the scores
scores = []
for id, e in evaluations.items():
    scores.append({"flow":e.flow_name, "score":e.value})

sorted_score = sorted(scores, key=lambda x: -x["score"])
fig, ax = pyplot.subplots(figsize=(8, 25))
sns.violinplot(ax=ax, x="score", y="flow", data=pd.DataFrame(sorted_score))

```



Easy benchmarking:

```
[25]: import openml as oml
      from sklearn import neighbors

      for task_id in [14951, 10103]:
          task = oml.tasks.get_task(task_id)
          data = oml.datasets.get_dataset(task.dataset_id)
          clf = neighbors.KNeighborsClassifier(n_neighbors=5)
          flow = oml.flows.sklearn_to_flow(clf)

          try:
              run = oml.runs.run_flow_on_task(task, flow)
              myrun = run.publish()
              print("kNN on %s: http://www.openml.org/r/%d" % (data.name, myrun.id))
          except oml.exceptions.PyOpenMLError as err:
              print("OpenML: {0}".format(err))
```

OpenML: Run already exists in server. Run id(s): {6068464}

OpenML: Run already exists in server. Run id(s): {6068467}

A Challenge

Try to build the best possible models on several OpenML tasks, and compare your results with the rest of the class, and learn from them. Some tasks you could try (or browse openml.org):

- EEG eye state: data_id:1471, task_id:14951
- Volcanoes on Venus: data_id:1527, task_id:10103
- Walking activity: data_id:1509, task_id: 9945, 150k instances
- Covertypes (Satellite): data_id:150, task_id: 218. 500k instances
- Higgs (Physics): data_id:23512, task_id:52950. 100k instances, missing values